

---

# Dynamic Model Construction for Efficient Classification

---

**Jaeeun Lee**

School of Computer Science  
University of Waterloo  
Waterloo, ON, N2L 3G1  
j474lee@uwaterloo.ca

## Abstract

One of the major drawbacks of neural network based classifier is that retraining is inevitable when target set of classes changes. For this reason, networks created for classification are often wide and deep to support all possible classes. This increases the amount of necessary computations and can lead to unpleasant user experience. In this work, I propose Composing algorithm which enables dynamic construction of a classifier using class-level transfer learning. Composing algorithm minimizes unnecessary computations but found to be less accurate once trained with cross entropy loss. I realize that sigmoid with binary cross entropy loss can minimize the accuracy decrease and include evaluation on how different loss function changes the behaviour of constructed model. From the experiments conducted on MNIST, keyword spotting, and CIFAR-100, it is found that sigmoid with binary cross entropy loss is more suitable for Composing algorithm but decrease in accuracy is inevitable as number of classes increases.

## 1 Introduction

Over the last decade, neural network has become *de facto* approach for numerous classification problems as it leads to high accuracy [1–3]. However, neural network based approaches require much larger computation and provide less flexible than preexisting techniques. When training a neural network based classifier, a set of target class must be provided in order to obtain a reliable classifier. The trained model then can be deployed and classify unseen data assuming that true class belongs to the target set which the model is trained on. However, in practice, this is not always the case and there exist two extreme cases where this setup falls apart.

The first case is when a set of true class contains only few classes from the target set. In this case, the trained model is considered to be an excessive representation of the true classifier and wastes computations as it calculates probabilities for unnecessary classes. This can be avoid when a set of true class is known prior to training, by setting a set of necessary classes as target classes.

The other case is when the true class does not exist in the target set. Unless the model is trained explicitly to classify such classes as unknown, the model will classify the unseen data to be one of the target classes and such misclassification can lead to a system failure. The ideal solution for this issues is to retrain a model with the new classes, minimizing the chance of misclassification.

However, training a neural network is a very expensive operation. It can take days to obtain a reliable classifier and this hinders the efficient management of a service. As a results, most of the academic works focus on minimizing resource usages of a network while preserving the high accuracy for every class. However, there exists an alternative solution to this problem: constructing a model dynamically adapting to the change in target set while minimizing decrease in accuracy and increase in resource usage. There are three conditions which the optimal solution must satisfy:

- 36 1. **Minimal accuracy degradation** : the difference in accuracy between a base model and  
37 corresponding composed model should be small
- 38 2. **Dynamic class addition and removal** : it must be easy to add and remove a class from the  
39 constructed model
- 40 3. **Efficient classification** : the constructed model should not require more computations than  
41 the base model

42 where base model refers to a model which trained explicitly to classify the same set of class which  
43 corresponding constructed model is trained to classify.

44 It is found that dynamic model construction is quite challenging as the optimal solution must con-  
45 sider relationship between each of the neurons and the output value of each class. In this paper, I  
46 present Composing algorithm which obtains such information by class-level transfer learning. As  
47 Composing algorithm involves mixing up the weights obtained from distinct models, I realize the  
48 limitation of standard cross entropy loss approach and show that sigmoid with binary cross entropy  
49 loss is more suitable for Composing algorithm. This has been demonstrated with a set of experiment  
50 as well. From the experiments conducted on MNIST, keyword spotting, and CIFAR-100, it is found  
51 that accuracy degradation is inevitable as number of classes increases but can be minimized when  
52 models are trained with sigmoid and binary cross entropy loss.

## 53 2 Related Works

54 Even though the three criteria for dynamic model construction are quite related, there exist distinct  
55 set of problems aiming to achieve each criteria. The three most relevant domains are: ensemble  
56 learning, multi-task learning, and transfer learning

### 57 2.1 Ensemble Learning

58 Ensemble learning is a common technique in the field of machine learning which achieves higher  
59 accuracy by combining outputs of multiple models. The most famous techniques include voting,  
60 weighting, bagging and boosting [4–6]. Even though ensemble learning is considered to be easy  
61 to implement, ensemble learning assumes that models are independent. As a result, most ensemble  
62 learning algorithms require each model to process the input data parallel violating the efficiency  
63 requirement of the dynamic model construction problem.

### 64 2.2 Multi-task Learning

65 On the other hand, multi-task learning takes the opposite approach; combine a set of networks to  
66 share the knowledge learned from each task. The key assumption is that if tasks are related, shar-  
67 ing knowledge throughout training will increase the performance of each network. Techniques for  
68 multi-task learning are often classified into two depending on the type of information being shared:  
69 parameter sharing and feature sharing [7–10]. Some of the techniques introduced for multi-task  
70 learning can be adapted to the dynamic model construction problem. However they fail to satisfy  
71 the efficiency requirements as extra layers are involved to share information while architectures of  
72 each model stay the same.

### 73 2.3 Transfer Learning

74 Transfer learning is inspired by the same assumption as multi-task learning; sharing knowledge  
75 among tasks can improve the performance. However, the key difference between two problems is  
76 that transfer learning use the same model architecture for multiple tasks. Transfer learning involves  
77 pre-training and fine-tuning. First, a model is pre-trained on a task and learns to select important  
78 features. Then the same model is fine-tuned for a target task as trained weights are adjusted to  
79 produce the best result for the target task [11]. It is found that transfer learning is very powerful  
80 as demonstrated by the wide range of applications [12–14]. Unlike two aforementioned domains,  
81 transfer learning does not add any computations to achieve better performance. However, knowledge  
82 sharing is mostly studied on task-level and limited work exists for class-level transfer learning.

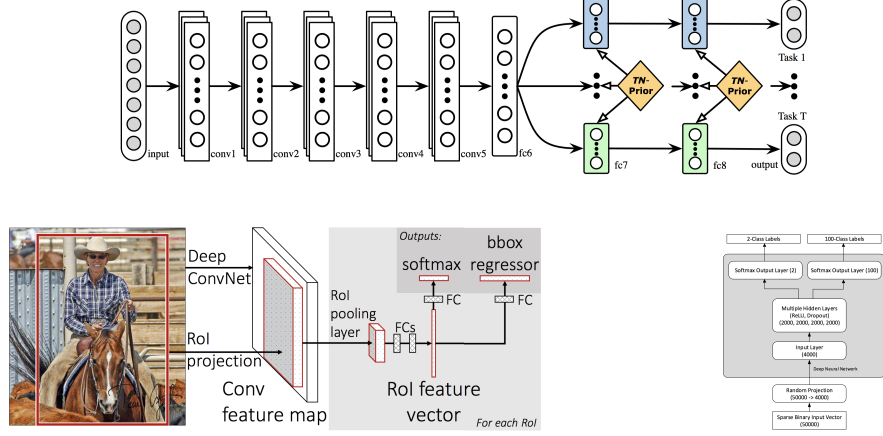


Figure 1: Model architectures proposed by [15] (top), [16] (bottom right), and [17] (bottom left); Their approach for multi-task learning is to assign distinct fully-connected layer for each task.

### 3 Composing Algorithm

In this section, I introduce Composing algorithm which enables dynamic model construction adapting to the change in target set. I also discuss the necessary conditions for preserving accuracy throughout Composing algorithm.

#### 3.1 Approach

One of the techniques proposed for multi-task learning is to share every layer while assigning distinct fully-connected layer for each task (see Figure 1). With such architecture, each task is trained independently while weights for upstream layers are shared among tasks. Once the training completes for the combined model, a task can be discarded from populating output by removing corresponding fully-connected layer.

Recall the criteria for dynamic model construction. Once the above multi-task learning approach is adjusted to support knowledge sharing among classes, dynamic addition and removal of a class is possible and the efficiency requirement can also be satisfied. Therefore, I introduce Composing algorithm which has following step (also see Figure 2):

1. Pre-train a model as if it is multi-classification problem
2. Freeze the model parameters
3. Construct a new dataset which labels one class as positive and the others as negative
4. Replace the last full-connected layer for two classes
5. Fine-tune the last layer using the new dataset
6. Retrieve the weights for positive class from the last fully-connected layer
7. Repeat step 4 ~ 7 for each class
8. For every combination of target classes, it is possible to construct a classifier by reconstructing the last layer with class-specific weights obtained from fine-tuning

In this work, a model obtained from pre-training as pre-trained model, models constructed from step 4 ~ 7 as fine-tuned models and a model constructed using this technique (obtained from the last step) as composed model.

Composing algorithm achieves dynamic addition and removal of a class by attaching and discarding corresponding fully-connected connections. With such flexibility, retraining a classifier is no longer necessary. Furthermore, Even though a class does not participate for pre-training, it is possible to add the class as long as it is fine-tuned using the same pre-trained model.




Figure 2: Visualization of Composing algorithm

Also, since composed models are built only using positive class connections of each fine-tuned model, the architecture is exactly same as pre-trained models and satisfy the efficiency criteria.

## 4 Accuracy Preservation

However, does Composing algorithm also guarantee the minimal accuracy degradation? To answer this question, we need to understand how different loss functions affect the accuracy of the composed model.

### 4.1 Limitations with cross entropy loss

State of the art loss function for multi-classification is cross entropy (CE) loss, which transforms output of a network by applying softmax and calculates negative log likelihood (NLL) as a measure of loss. Softmax and NLL loss are defined as following:

$$NegativeLogLikelihood(y, t) = -\frac{1}{N} \sum_{i=1}^N [t_i \cdot \log y_i]$$

$$y_i = Softmax(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

where  $x$  is the output of a network and  $t$  is a target, one hot encoded vector. The definition of softmax can be summarized as calculating normalized logits using the network output. Therefore,  $y$  has values between zero and one and they must sum up to one. Since multi-class classifier rely on the mutually exclusive assumption among the classes, CE loss is found to be the most powerful as it promotes the positive class while suppressing the negative classes.

However, such assumption can lead to unpredictable behaviour with Composing algorithm. When CE loss is used for Composing algorithm, the losses calculated throughout fine-tuning involve outputs from corresponding negative class. However, for a composed model, probability for each class is computed using the class-specific weights obtained from independently fine-tuned models and a class with the highest probability is selected to be the final prediction.

For example, let us say there are three classes: A, B, and C. For fine-tuning, each dataset consists of two classes: a positive class and a negative class which represents all the other classes. For simplicity, positive weights are referred with lower case alphabets (a, b, and c) and negative weights with

136 lower case alphabets with prime (a', b', and c'). Throughout fine-tuning process, loss is calculated  
 137 in pairs as following: a – a', b – b', c – c'. However, for a combined model, last layer is constructed  
 138 with weights a, b and c. As a results, there is no guarantee that the class with the highest output is  
 139 in fact the class with the highest probability.

## 140 4.2 Binary cross entropy with sigmoid

141 Given that CE loss does not guarantee the same accuracy due to mutually exclusive. Therefore, I  
 142 analyze sigmoid with binary cross entropy (BCE) loss.

$$BinaryCrossEntropy(y, t) = -\frac{1}{N} \sum_{i=1}^N [t_i \cdot \log y_i + (1 - t_i) \cdot \log(1 - t_i)]$$

$$y_i = Sigmoid(y_i) = \frac{1}{1 + e^{-x_i}}$$

143 Unlike CE loss, both sigmoid and BCE loss treat each output independently. In other words, weights  
 144 for the positive class in each fine-tuned model no longer depend on the negative class. This indicates  
 145 that the output of composed models are more reliable.

146 In multi-label classification, the same issue has been raised with CE loss [18]. It is found that the  
 147 independence guarantee provided by sigmoid and BCE loss is crucial for multi-label classification  
 148 and enables successful training of a classifier.

## 149 5 Experiments

150 In order to understand the severity of accuracy decrease, I have implemented Composing algorithm  
 151 on MNIST, Keyword Spotting, and CIFAR-100 using PyTorch and available on github<sup>1</sup>.

152 For each dataset, Composing algorithm is evaluate with three loss functions. The first loss function  
 153 is CE loss. Since PyTorch NLL loss implementation expects log probability, log is applied after  
 154 softmax but this does not affect the analysis. Next, I use sigmoid with BCE loss as it is found to  
 155 be more suitable than CE loss. Last loss function is softmax with BCE loss. This setting is known  
 156 to be unstable because BCE loss assumes the independence among classes while softmax does not.  
 157 In fact, I have observed the training collapse at some point. However, as I report accuracy from the  
 158 best model, I found the results from this setting still valid and meaningful. This combination should  
 159 allow me to understand how crucial sigmoid is for sigmoid with BCE loss as it simply replaces  
 160 sigmoid with softmax.

161 In the following sections, I report accuracy of every model created throughout Composing algorithm:  
 162 pre-trained, fine-tuned, and composed models. My main goal of this experiments is to understand  
 163 how each loss function affects the stability of Composing algorithm. In order to analyze such rela-  
 164 tionship, I compare pre-trained model accuracy against composed model accuracy.

165 Furthermore, I report accuracy of composed model varying number of classes. This reveals re-  
 166 lationship between number of classes and the performance of composed model. Since fine-tuned  
 167 accuracies varies a lot depending on the class, I report an average from 10 iterations with random  
 168 selection on the class to add next.

### 169 5.1 MNIST

170 MNIST is a standard benchmark for classification which comprises images of handwritten digits [1].  
 171 Among the wide range of model architectures proposed for this problem, LeNet-5 is selected  
 172 for this experiment. LeNet-5 is constructed with two convolutional, one dropout and two fully  
 173 connected layers [19]. The original implementation of LeNet-5 has 10 and 20 channels for the  
 174 first two convolutional layers and produces accuracy of 98% on MNIST dataset. Since accuracy is

<sup>1</sup><https://github.com/ljj7975/composable-model-exp>

Loss function	Pre-trained	Fine tuned avg (min ~ max)	Composed	Relative decrease
LogSoftmax + NLL	95.8	98.02 (96.29 ~ 99.24)	85.74	10.50
Softmax + BCE	94.71	97.33 (95.08 ~ 99.06)	77.80	17.85
Sigmoid + BCE	95.49	98.07 (96.74 ~ 99.19)	95.30	0.20

Table 1: Average accuracy of base, fine-tuned, and composed model for MNIST (%). Relative decrease is calculated with respect to pre-trained model.




Figure 3: Change in composed model accuracy with respect to number of classes for MNIST

the prior measure of comparison in this experiments, such a high accuracy might lead to difficult analysis. Therefore, the network is limited with 5 channels for both convolutional layers.

Both pre-trained model and fine-tuned models are trained using Adam optimizer with learning rate of 0.0001. From 50 experiments, it is found that all three loss function lead to convergence in 5 epochs with 95% accuracy for pretraining and 98% accuracy for fine-tuning (see Table 1).

Figure 3 summarizes how the accuracy changes for each composed model as number of classes increases. No matter which loss function is used, accuracy of composed model decreases as more classes contribute. However, models with softmax based loss show greater rate of decrease than a model with sigmoid based loss. With CE loss, the composed model for all 10 classes show accuracy of 85.86%. This is relative decrease of 10.5% from the pre-trained model. Composing algorithm using softmax with BCE loss shows the worst performance. The average composed model accuracy is 78.50% which is 17.85% relative decrease. As shown in the previous section, sigmoid with BCE loss introduces the least accuracy degradation and achieves accuracy of 95.29% which is very similar to the accuracy of pre-trained model.

## 5.2 Keyword Spotting

To understand the universality of Composing algorithm, I extend this idea to keyword spotting where the input data is audio. The goal of Keyword Spotting (KWS) is to detect an audio of pre-trained keywords, such as “Hey Siri”. Since the success of deep neural network approach by [2], neural network has become the standard approach for KWS. In this experiment, I implement `res15-narrow` introduced by [20], which achieves 94% accuracy for 12 keywords on Google’s Speech Commands Dataset [21]. `res15-narrow` comprises 6 residual blocks with 19 feature maps where each residual block is composed of bias-free convolutional and batch normalization layer.

Common KWS experiments on Google’s Speech Commands Dataset involve only 12 keywords. However, since I am interested in evaluating stability of Composing algorithm with larger number of classes, all 30 keywords are used for this experiment. Following the standard feature extraction for audio data, I first construct Forty-dimensional Mel-Frequency Cepstrum Coefficient (MFCC) frames and stack them using 30ms windows with a 10ms shift. Since the dataset consists of one-second long utterances of each word, the final input has size of  $101 \times 40$ .

Throughout the 10 experiments, stochastic gradient descent is used for both pre-training and fine-tuning. Training starts with learning rate of 0.1 and achieves higher accuracy as it decreases the learning rate to 0.001 by factor of ten. Base models are trained for 30 epochs with learning rate

Loss function	Pre-trained	Fine tuned avg (min ~ max)	Composed	Relative decrease
LogSoftmax + NLL	93.09	95.32 (92.57 ~ 97.59)	90.13	3.18
Softmax + BCE	90.94	91.79 (89.64 ~ 94.80)	86.78	4.57
Sigmoid + BCE	89.62	91.31 (88.73 ~ 93.91)	88.33	1.44

Table 2: Average accuracy of base, fine-tuned, and composed model for KWS (%). Relative decrease is calculated with respect to pre-trained model.




Figure 4: Change in composed model accuracy with respect to number of classes for KWS

decrease at 10 and 20th epochs and fine-tuned models are trained for 10 epochs with decrease at 4 and 7th epochs.

Unlike MNIST, it is found that all three composed models show reasonably good accuracy. First, CE loss achieves the best accuracy with `res15-narrow`; 93.09% from pre-training and average accuracy of 95.32% with fine-tuning. Softmax with BCE loss achieves 90.94% accuracy with pre-trained model and average accuracy of 91.79% with fine-tuned models. Sigmoid with BCE loss leads to the least accuracy of 89.62% from pre-training and 91.31% from fine-tuning.

However, Figure 4 shows that sigmoid with BCE loss is the most reliable loss function as it shows the least relative decrease of 1.44%. CE loss and softmax with BCE loss show greater rate of relative decrease, 3.18% and 4.57% respectively. As observed with MNIST, decrease in accuracy is also found with KWS as number of classes increases.

### 5.3 CIFAR-100

CIFAR is a collection of tiny coloured images from the web [3]. There exist two variations for CIFAR differing number of classes: CIFAR-10 and CIFAR-100. The following experiment is constructed with CIFAR-100 which constitutes 600 images of 100 classes.

For this experiment, I have implemented DenseNet, a state of the art model for CIFAR dataset [22]. Building upon a network architecture with residual connection, the feature maps of all preceding layers are used as inputs for each layer. The network has three dense blocks with transition layers between which changes the feature-map sizes by convolution and pooling. The original implementation achieves accuracy of 80% with 300 epochs of stochastic gradient descent. Learning rate must decrease throughout the training from 0.1 to 0.001 by factor of ten.

Table 3 summarizes the results from 5 experiments on CIFAR-100. The pre-training is achieved with 200 epochs and this leads to accuracy of 69.95% for CE loss, 64.23% for softmax with BCE loss, and 64.72% for sigmoid with BCE loss. Fine-tuned models are trained for 100 epochs and each loss function converges to average accuracy of 86.12%, 88.63%, and 87.79% respectively.

Figure 5 shows how accuracy of composed model changes as number of classes increases. It is found that CIFAR-100 introduces greater rate of decrease than MNIST and KWS. I believe this is due to the fact that CIFAR-100 involves much larger number of classes. Again, limitation of CE loss is clear as it leads to 24.60% relative decrease from the base model. Softmax with BCE loss shows 18.98% relative decrease while sigmoid with BCE loss only show the least decrease of 11.28%.

Loss function	Pre-trained	Fine tuned avg (min ~ max)	Composed	Relative decrease
LogSoftmax + NLL	69.95	86.12 (71.00 ~ 96.00)	52.74	24.60
Softmax + BCE	64.23	88.63 (79.50 ~ 97.50)	52.04	18.98
Sigmoid + BCE	64.72	87.79 (77.50 ~ 96.00)	57.42	11.28

Table 3: Average accuracy of base, fine-tuned, and composed model for CIFAR-100 (%). Relative decrease is calculated with respect to pre-trained model.




Figure 5: Change in composed model accuracy with respect to number of classes for CIFAR-100

## 6 Discussion

Throughout all three experiments, strong correlation is observed between number of classes and composed model accuracy. This leads to severe accuracy degradation for CIFAR-100 as it involves 100 classes. Therefore a composed model which does not suffer from increase in number of classes would be preferred when there exist large number of classes. I believe the first step is to conduct further experiments using other type of loss functions such as KL divergence and MSE, understanding how these loss functions affects the performance.

Next, when I introduce Composing algorithm, I only interact with the last fully-connected layer. Therefore, one might believe that the algorithm only works with such network. However, this algorithm can be extended to other network as long as correct layer is selected for fine-tuning. For example, some networks apply global averaging layer instead of fully-connected layer to minimize computation. Since global averaging does not involve any parameter, the last valid layer for fine-tuning is penultimate layer. In such networks, penultimate layer is generally a convolutional layer. Therefore, as long as fine-tuning leads to reasonable accuracy and weights are loaded correctly for the new penultimate layer, the same Composing algorithm should work on these networks.

Lastly, I have shown that Composing algorithm saves computation as it does not calculate probability for unnecessary classes. However, as model involves more and more layers, such savings may not add much benefit since Composing algorithm only saves computations on the last layer. In order for this work to be meaningful, it is necessary to extend this idea for upstream layers and achieve greater rate of savings on computation.

## 7 Conclusion

Realizing the limited flexibility of a classifier, I present Composing algorithm which enables dynamic construction of a model adapting to the change in target classes. It supports addition and removal of a class on the fly, achieving computational efficiency without any further training. I show how CE loss can lead to severe accuracy degradation with Composing algorithm and realize that sigmoid with BCE loss guarantees better accuracy. It is also demonstrated empirically using MNIST, keyword spotting and CIFAR-100 as I analyze how different loss functions affect the performance of composed model. Unfortunately, it is also found that this algorithm suffers from greater rate of accuracy decrease as number of classes increases, which can be crucial in some cases.



## References

- [1] Yann LeCun, Léon Bottou, Yoshua Bengio, Patrick Haffner, et al. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- [2] Guoguo Chen, Carolina Parada, and Georg Heigold. Small-footprint keyword spotting using deep neural networks. In *2014 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 4087–4091. IEEE, 2014.
- [3] Alex Krizhevsky and Geoffrey Hinton. Learning multiple layers of features from tiny images. Technical report, Citeseer, 2009.
- [4] Thomas G Dietterich. Ensemble methods in machine learning. In *International workshop on multiple classifier systems*, pages 1–15. Springer, 2000.
- [5] Leo Breiman. Bagging predictors. *Machine learning*, 24(2):123–140, 1996.
- [6] Yoav Freund, Robert E Schapire, et al. Experiments with a new boosting algorithm. In *icml*, volume 96, pages 148–156. Citeseer, 1996.
- [7] Sebastian Ruder. An overview of multi-task learning in deep neural networks. *arXiv preprint arXiv:1706.05098*, 2017.
- [8] Rich Caruana. Multitask learning: A knowledge-based source of inductive bias. In *ICML*, 1993.
- [9] Long Duong, Trevor Cohn, Steven Bird, and Paul Cook. Low resource dependency parsing: Cross-lingual parameter sharing in a neural network parser. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, volume 2, pages 845–850, 2015.
- [10] Yongxi Lu, Abhishek Kumar, Shuangfei Zhai, Yu Cheng, Tara Javidi, and Rogerio Feris. Fully-adaptive feature sharing in multi-task networks with applications in person attribute classification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 5334–5343, 2017.
- [11] Jason Yosinski, Jeff Clune, Yoshua Bengio, and Hod Lipson. How transferable are features in deep neural networks? In *Advances in neural information processing systems*, pages 3320–3328, 2014.
- [12] Rajat Raina, Alexis Battle, Honglak Lee, Benjamin Packer, and Andrew Y Ng. Self-taught learning: transfer learning from unlabeled data. In *Proceedings of the 24th international conference on Machine learning*, pages 759–766. ACM, 2007.
- [13] Toby Marshall Egan, Baiyin Yang, and Kenneth R Bartlett. The effects of organizational learning culture and job satisfaction on motivation to transfer learning and turnover intention. *Human resource development quarterly*, 15(3):279–301, 2004.
- [14] Xavier Glorot, Antoine Bordes, and Yoshua Bengio. Domain adaptation for large-scale sentiment classification: A deep learning approach. In *Proceedings of the 28th international conference on machine learning (ICML-11)*, pages 513–520, 2011.
- [15] Mingsheng Long, Zhangjie Cao, Jianmin Wang, and S Yu Philip. Learning multiple tasks with multilinear relationship networks. In *Advances in Neural Information Processing Systems*, pages 1594–1603, 2017.
- [16] Wenyi Huang and Jack W Stokes. Mtnet: a multi-task neural network for dynamic malware classification. In *International Conference on Detection of Intrusions and Malware, and Vulnerability Assessment*, pages 399–418. Springer, 2016.
- [17] Ross Girshick. Fast r-cnn. In *Proceedings of the IEEE international conference on computer vision*, pages 1440–1448, 2015.
- [18] Jingzhou Liu, Wei-Cheng Chang, Yuexin Wu, and Yiming Yang. Deep learning for extreme multi-label text classification. In *Proceedings of the 40th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 115–124. ACM, 2017.
- [19] Yann LeCun et al. Lenet-5, convolutional neural networks. URL: <http://yann.lecun.com/exdb/lenet>, 20, 2015.

- 316 [20] Raphael Tang and Jimmy Lin. Deep residual learning for small-footprint keyword spotting. In  
317 *2018 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*,  
318 pages 5484–5488. IEEE, 2018.
- 319 [21] Pete Warden. Launching the speech commands dataset. [https://research.  
320 googleblog.com/2017/08/launching-speech-commands-dataset.html](https://research.googleblog.com/2017/08/launching-speech-commands-dataset.html),  
321 2017.
- 322 [22] Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q Weinberger. Densely con-  
323 nected convolutional networks. In *Proceedings of the IEEE conference on computer vision and  
324 pattern recognition*, pages 4700–4708, 2017.